

# Star Schema Benchmark für SAP HANA

Jan Hofmeier, Kristina Albrecht, Lion Scherer, Marius Jochheim

1. SAP HANA
2. Star Schema Benchmark
3. Benchmarks
4. Cube Präsentation

---

## Gliederung

# SAP HANA

---

High Performance Analytic  
Appliance

---

Hauptspeicher statt Festplatte

---

Spaltenorientiert (Row store  
möglich)

---

OLAP + OLTP

---

Entwicklungsplattform

---

# Speicherzugriffszeiten

Speicherkomponenten in der Systemarchitektur	Größenordnung der Zugriffszeit
Zugriff auf CPU L1-/L2-/ L3 Cache	0,5 / 7,0 / 15 ns
Zugriff auf Hauptspeicher	100 ns
Zugriff auf Solid-State-Festplatte (SSD)	150.000 ns
Festplattenzugriff	10.000.000 ns

Bildquelle: In-Memory-Datenbank SAP HANA,  
Peter Preuss, Springer Gabler 2017

# In Memory Vorteile

---

Komplexe Abfragen und Datenbankoperationen  
mit sehr hohem Durchsatz ausführbar

---

Kein ETL Prozess

---

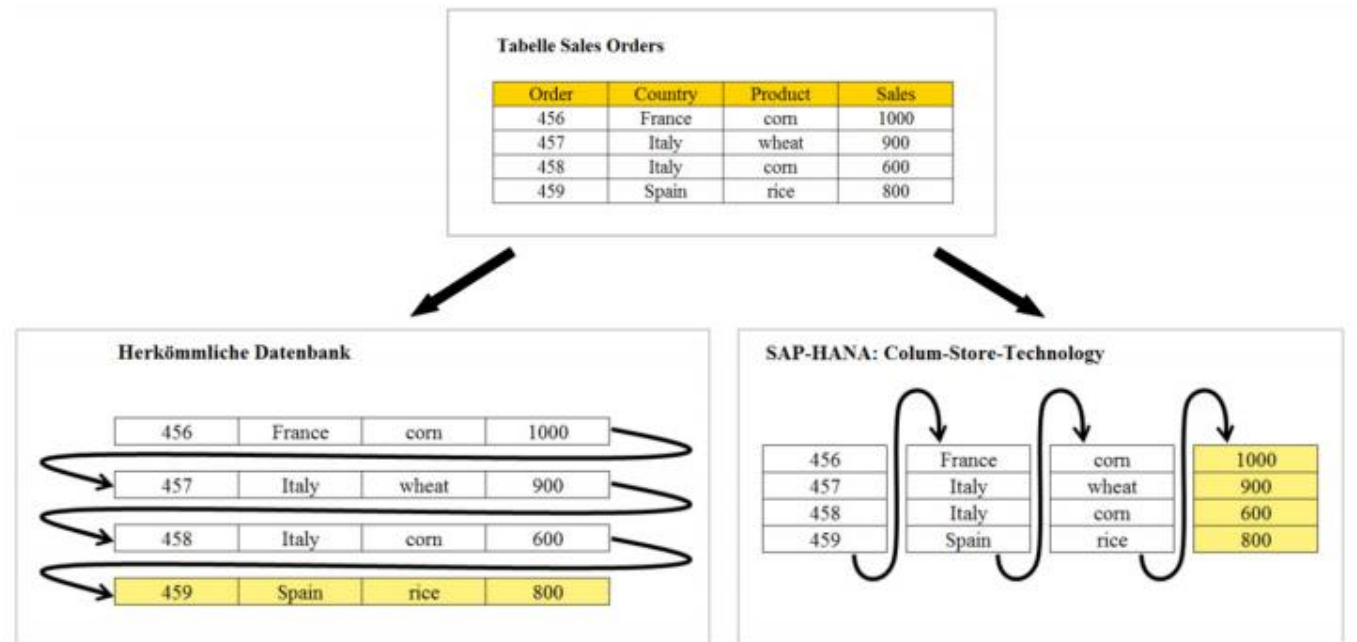
Echtzeit Analysen

---

Datenzugriffszeit wird verkürzt

---

# Spaltenbasiert vs Zeilenbasiert



Bildquelle: In-Memory-Datenbank SAP HANA, Peter Preuss,  
Springer Gabler 2017

# Compression

Dictionary Compression

Prefix encoding

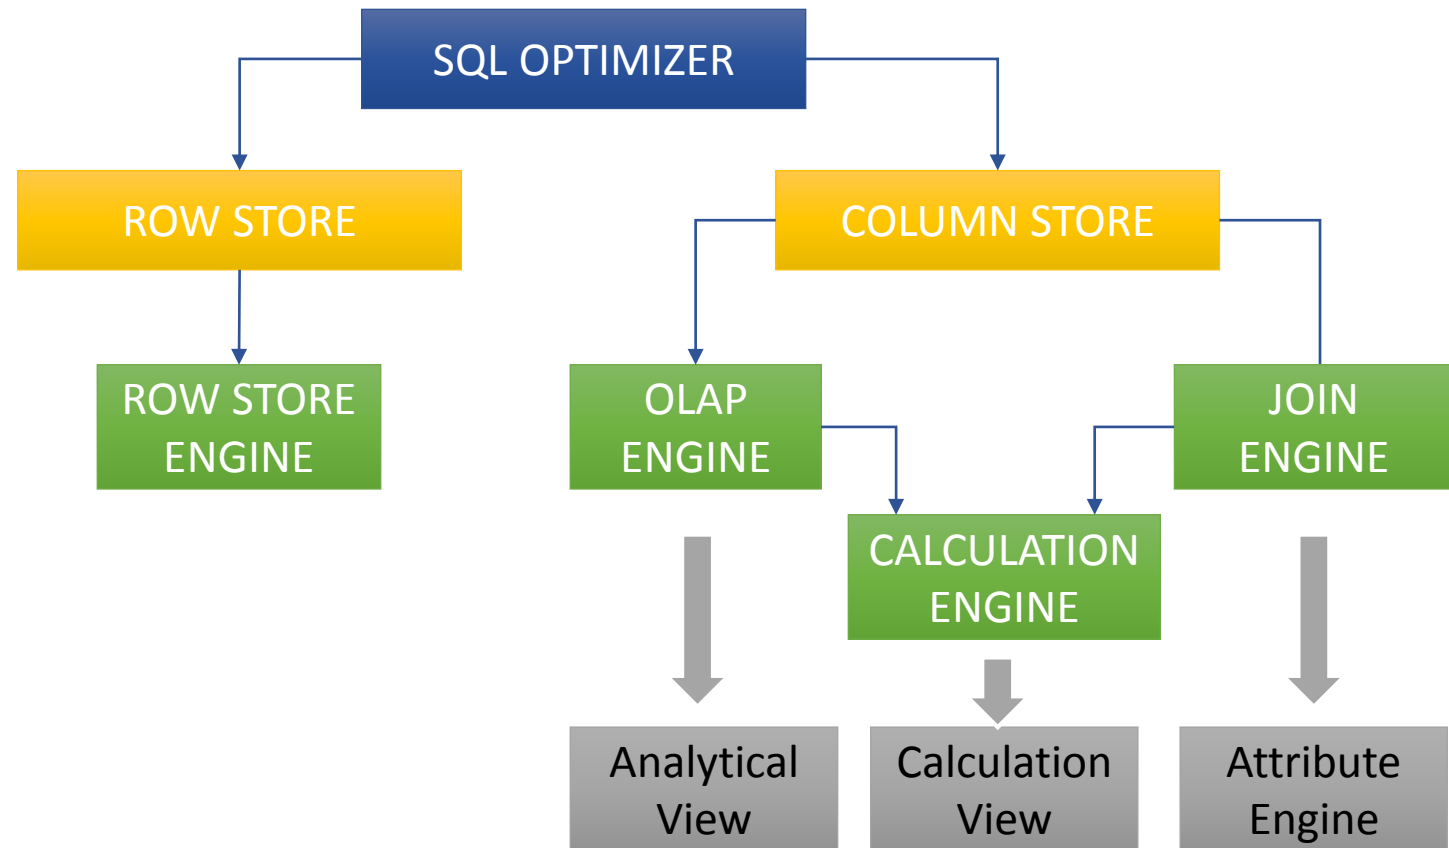
Run length encoding

Cluster encoding

Sparse encoding

Indirect encoding


# Architektur



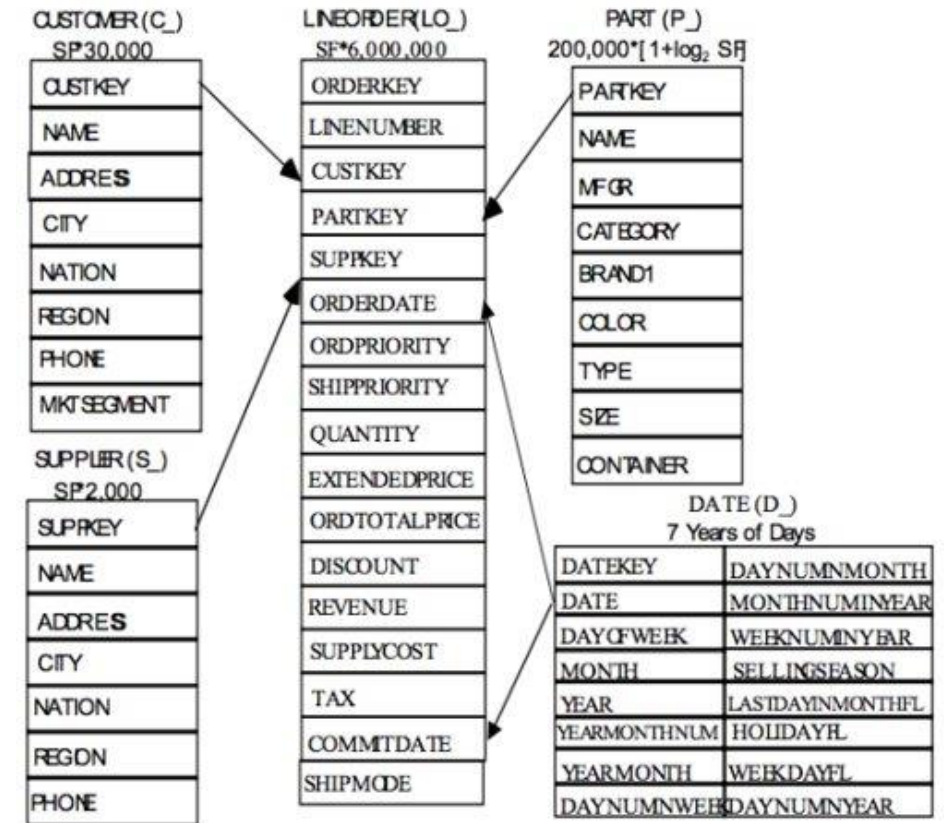
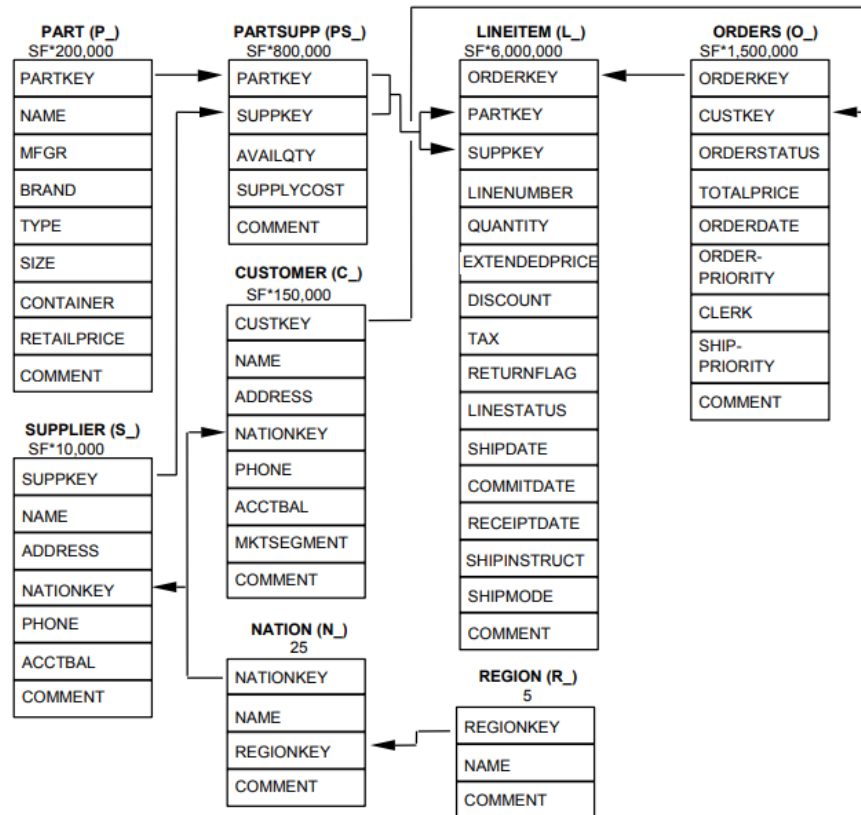


# Star Schema Benchmark

9



- Misst die OLAP Performance von Datenbank Produkten
- Basiert auf dem TPC-H Benchmark
- Nutzt ein Star Schema Data Mart mit großer Faktentabelle (LINEORDER) und mehreren Dimensionstabellen



# Vergleich von TPC-H und SSF

# Group 1

```
SELECT SUM(lo_extendedprice*lo_discount)
AS revenue
FROM lineorder JOIN dim_date
ON lo_orderdatekey = d_datekey
WHERE d_year = 1993
AND lo_discount BETWEEN 1 AND 3
AND lo_quantity < 25;
```

## Group 2

```
SELECT SUM(lo_revenue), d_year, p_brand
FROM lineorder
JOIN dim_date ON lo_orderdatekey = d_datekey
JOIN part ON lo_partkey = p_partkey
JOIN supplier ON lo_suppkey = s_suppkey
WHERE p_brand= 'MFGR#2239'
AND s_region = 'EUROPE'
GROUP BY d_year, p_brand
ORDER BY d_year, p_brand;
```

## Group 3

```
SELECT c_city, s_city, d_year, SUM(lo_revenue) AS revenue
FROM customer
JOIN lineorder ON lo_custkey = c_customerkey
JOIN supplier ON lo_suppkey = s_suppkey
JOIN dim_date ON lo_orderdatekey = d_datekey
WHERE (c_city='UNITED KI1' OR c_city='UNITED KI5')
AND (s_city='UNITED KI1' OR s_city='UNITED KI5')
AND d_year >= 1992
AND d_year <= 1997
GROUP BY c_city, s_city, d_year
ORDER BY d_year ASC, revenue DESC;
```

# Group 4

```
SELECT d_year, s_city, p_brand,  
SUM(lo_revenue - lo_supplycost) AS profit  
FROM lineorder  
JOIN dim_date ON lo_orderdatekey = d_datekey  
JOIN customer ON lo_custkey = c_customerkey  
JOIN supplier ON lo_suppkey = s_suppkey  
JOIN part ON lo_partkey = p_partkey  
WHERE s_nation = 'UNITED STATES'  
AND (d_year = 1997 OR d_year = 1998)  
AND p_category = 'MFGR#14'  
GROUP BY d_year, s_city, p_brand  
ORDER BY d_year, s_city, p_brand;
```

# HANA Performance Benchmark



# HANA Performance Benchmark

## Test Setup

VM Image von SAP

---

CPU: Intel(R) Core(TM) i7-6820HQ CPU @ 2.70GHz

---

CPU Kerne: 4 (x2 Theads)

---

RAM: 16GB

---

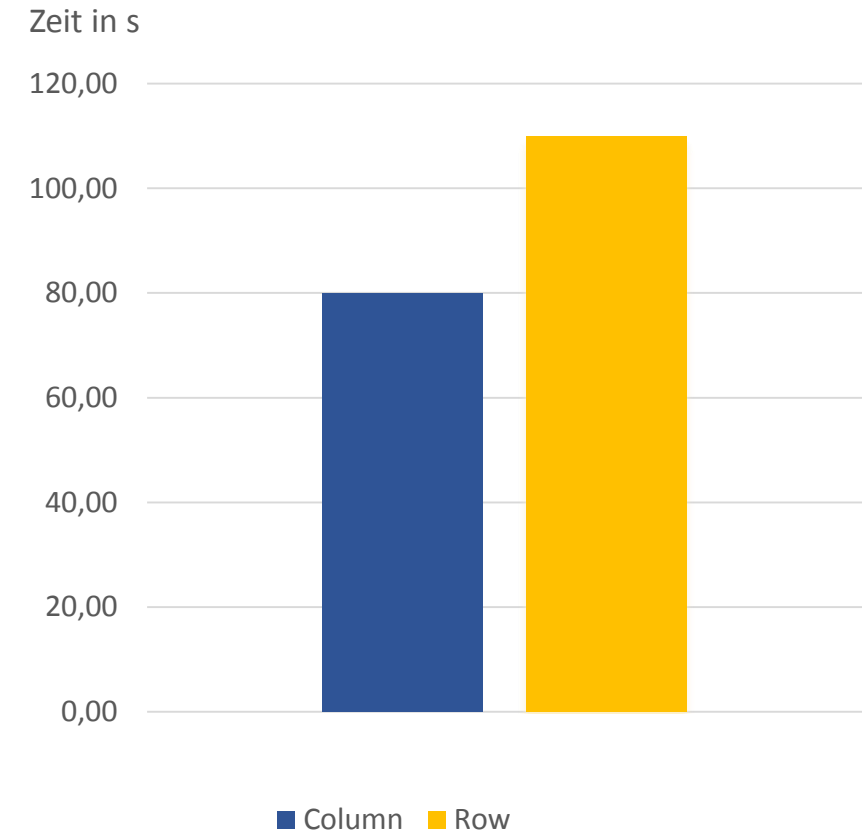
Storage: SSD



# HANA Performance Benchmark

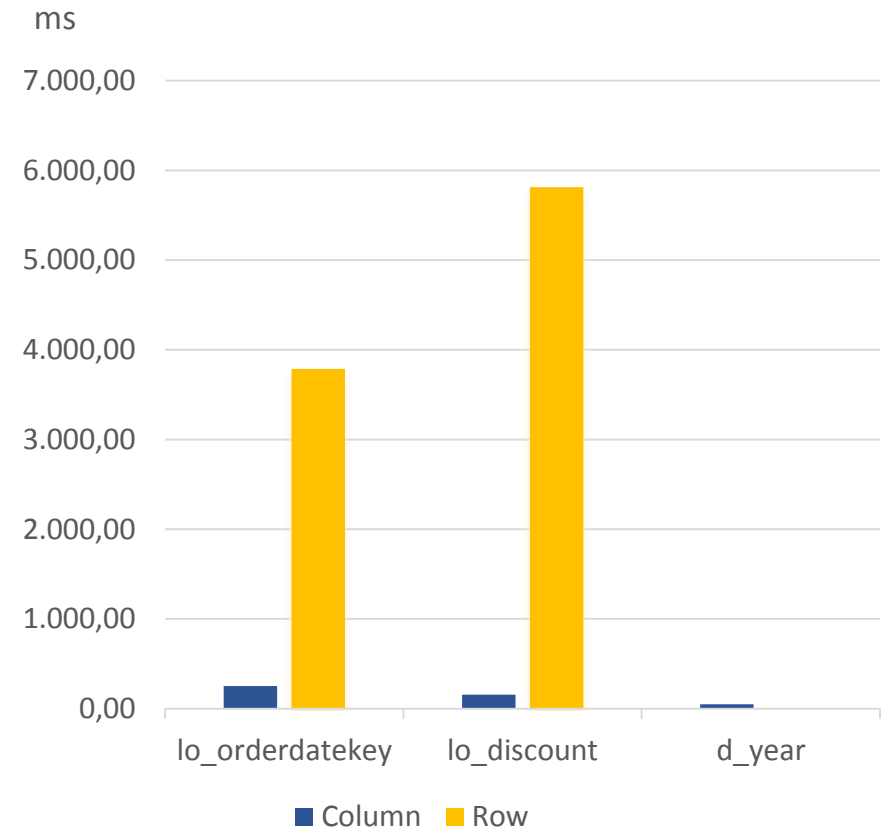
## Ladezeiten von Tabellen

**27 % weniger Zeit**



# HANA Performance Benchmark

## Ladezeiten von Indizes



# HANA Performance Benchmark

Vorgehensweise

**Row vs. Column Store**

---

**Indizes**

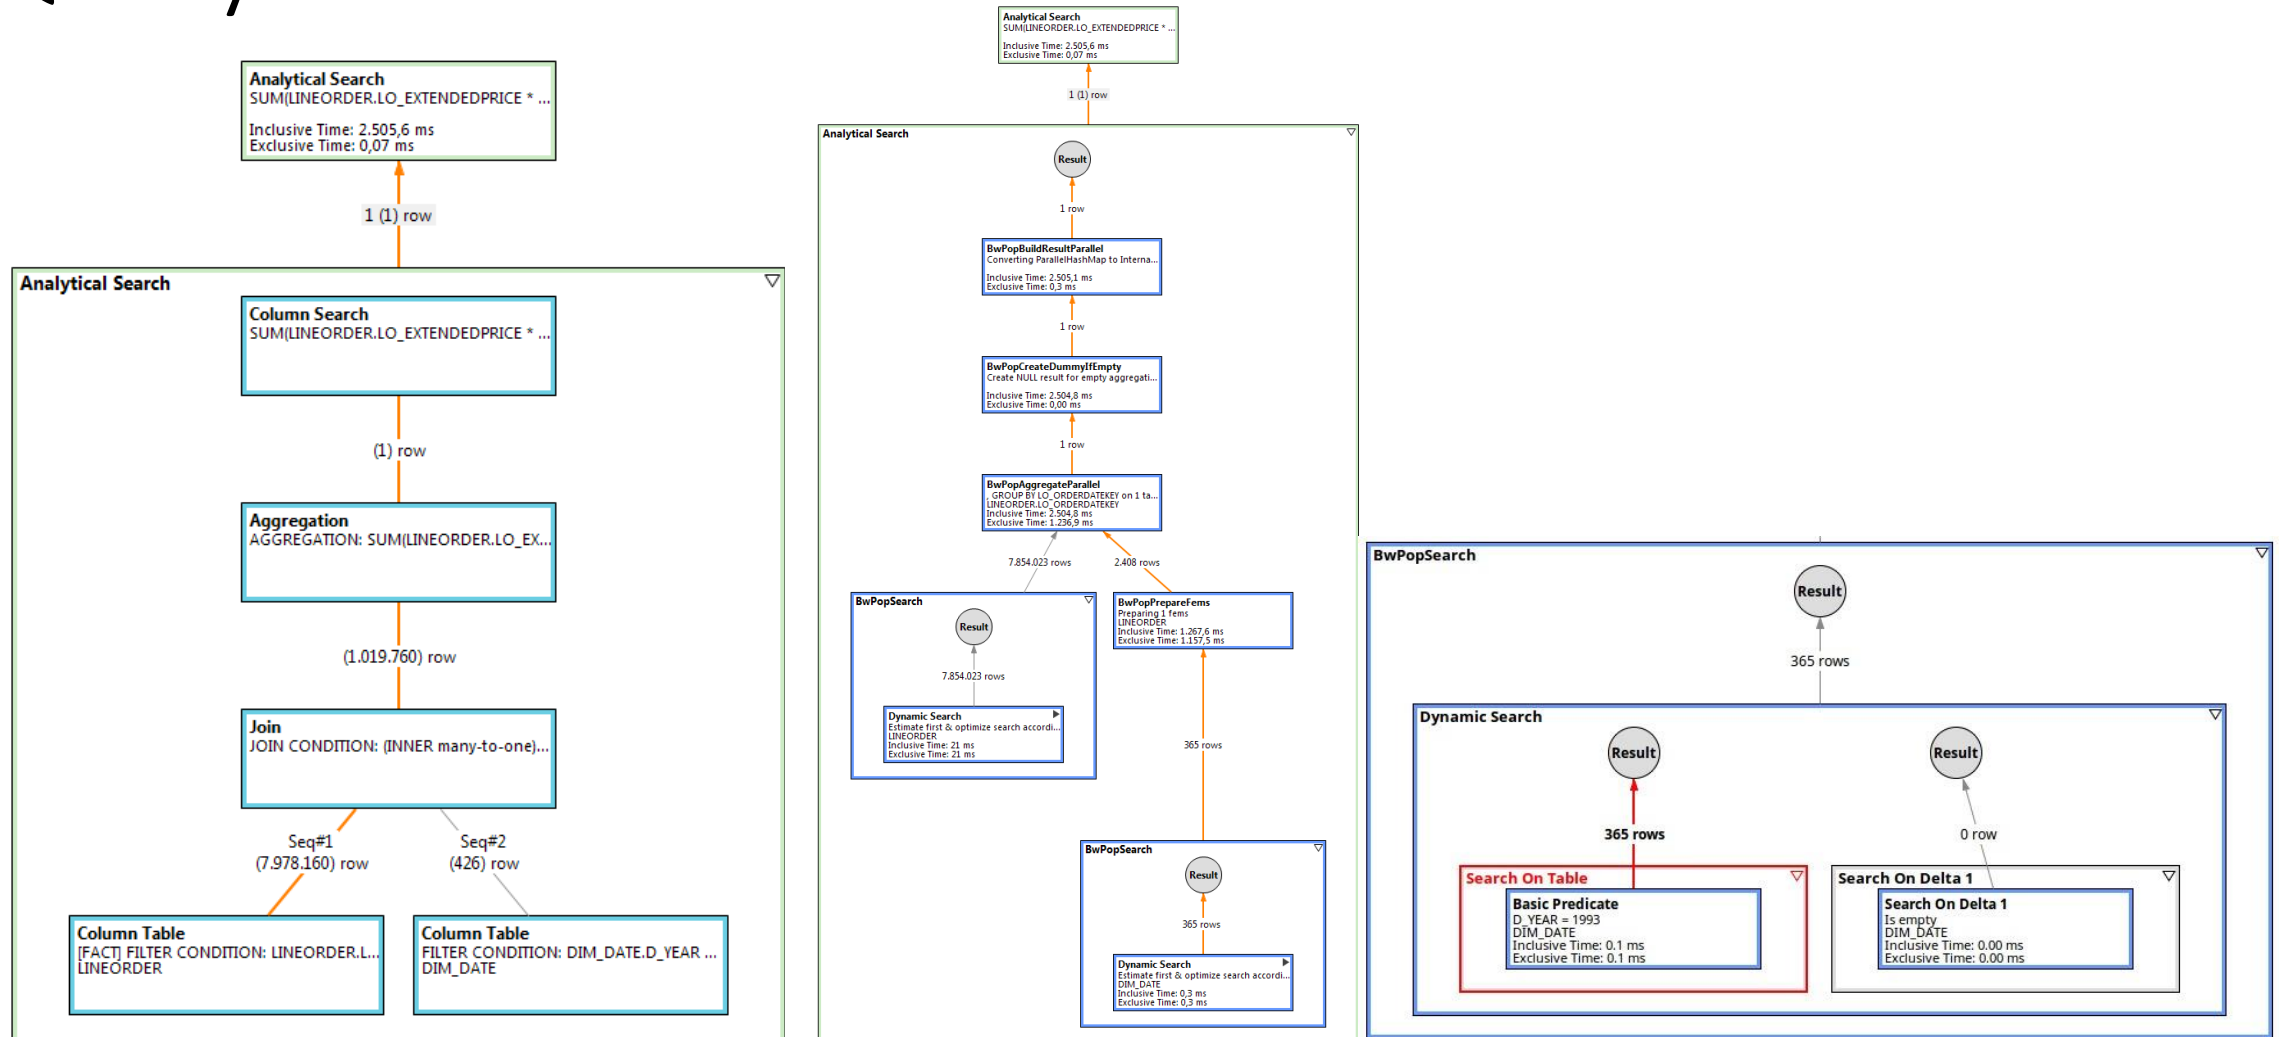
---

**Hints**


---

Foreign Key (FK)	+ Faktentabelle (FT)	+ Nur Dimensionen (DimOnly)	+ Restriktive Indizes auf Dimensionen (RestrDim)	Keine Indizes (None)
lo_custkey	lo_quantity	c_region	c_city	
lo_suppkey	lo_extendedprice	c_mktsegment	p_brand	
lo_partkey	lo_discount	p_mfgr	s_city	
lo_orderdatekey		p_category	d_yearmonthnum	
lo_commitdatekey		s_nation	d_yearmonth	
		s_region	d_daynuminyear	
		d_year		


# Query Execution Plans



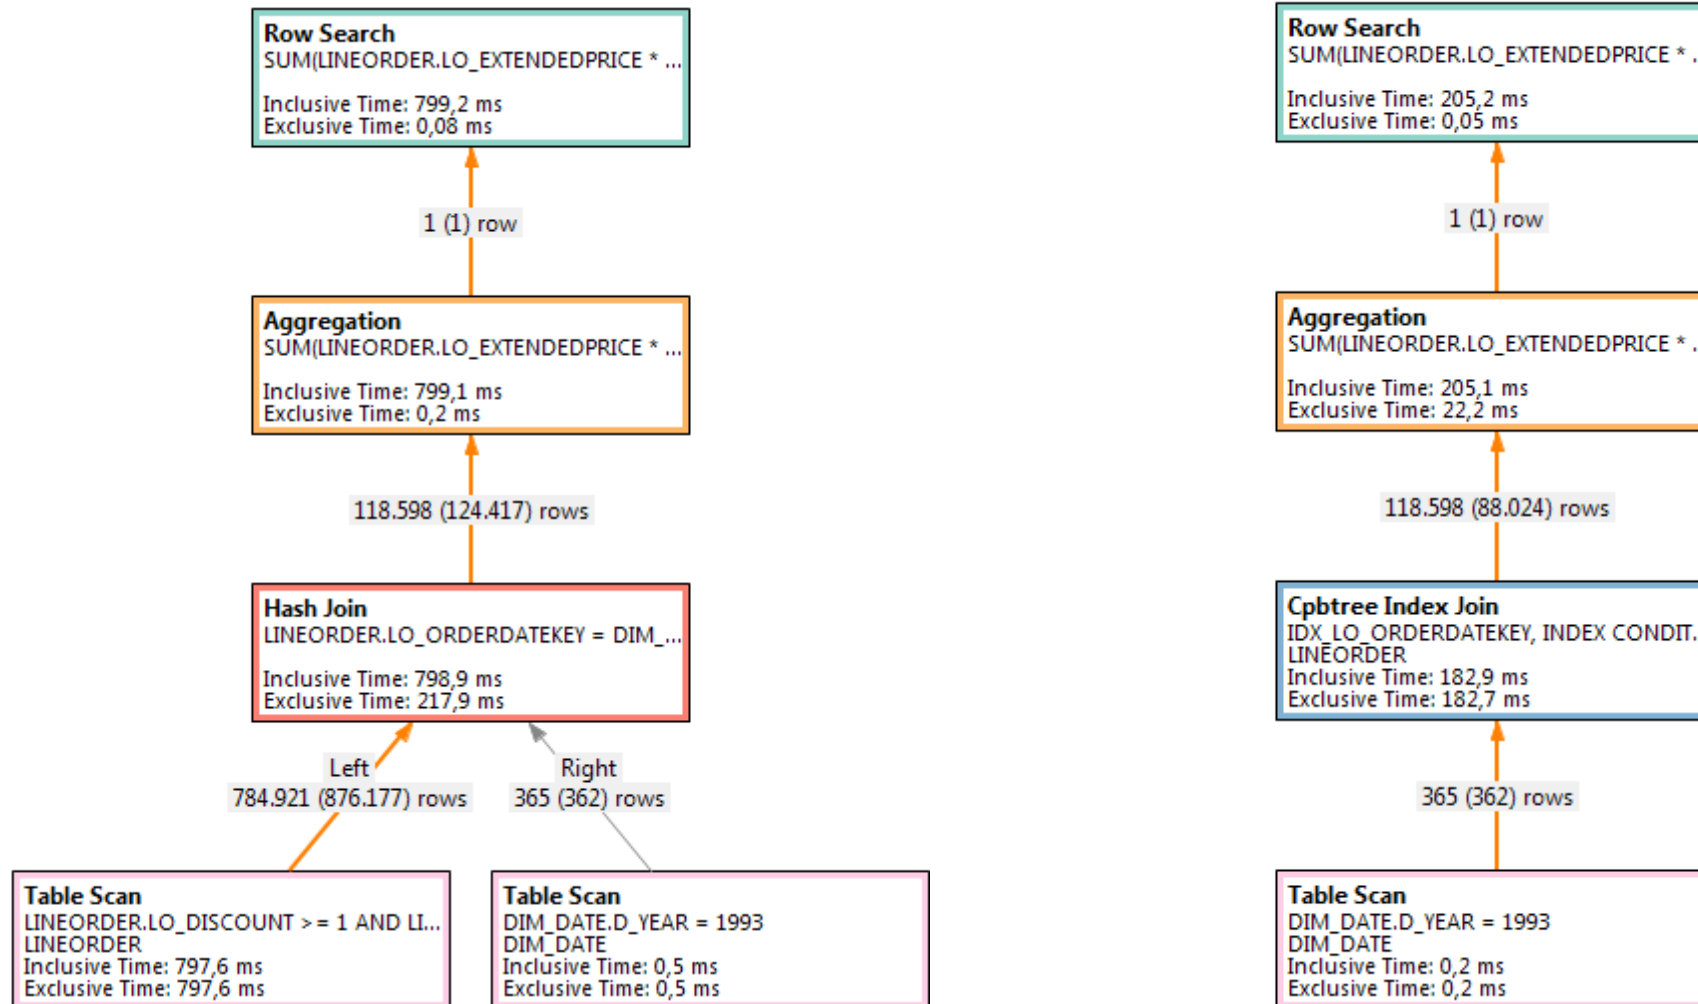
# Column vs. Row Store



# Column vs. Row Store

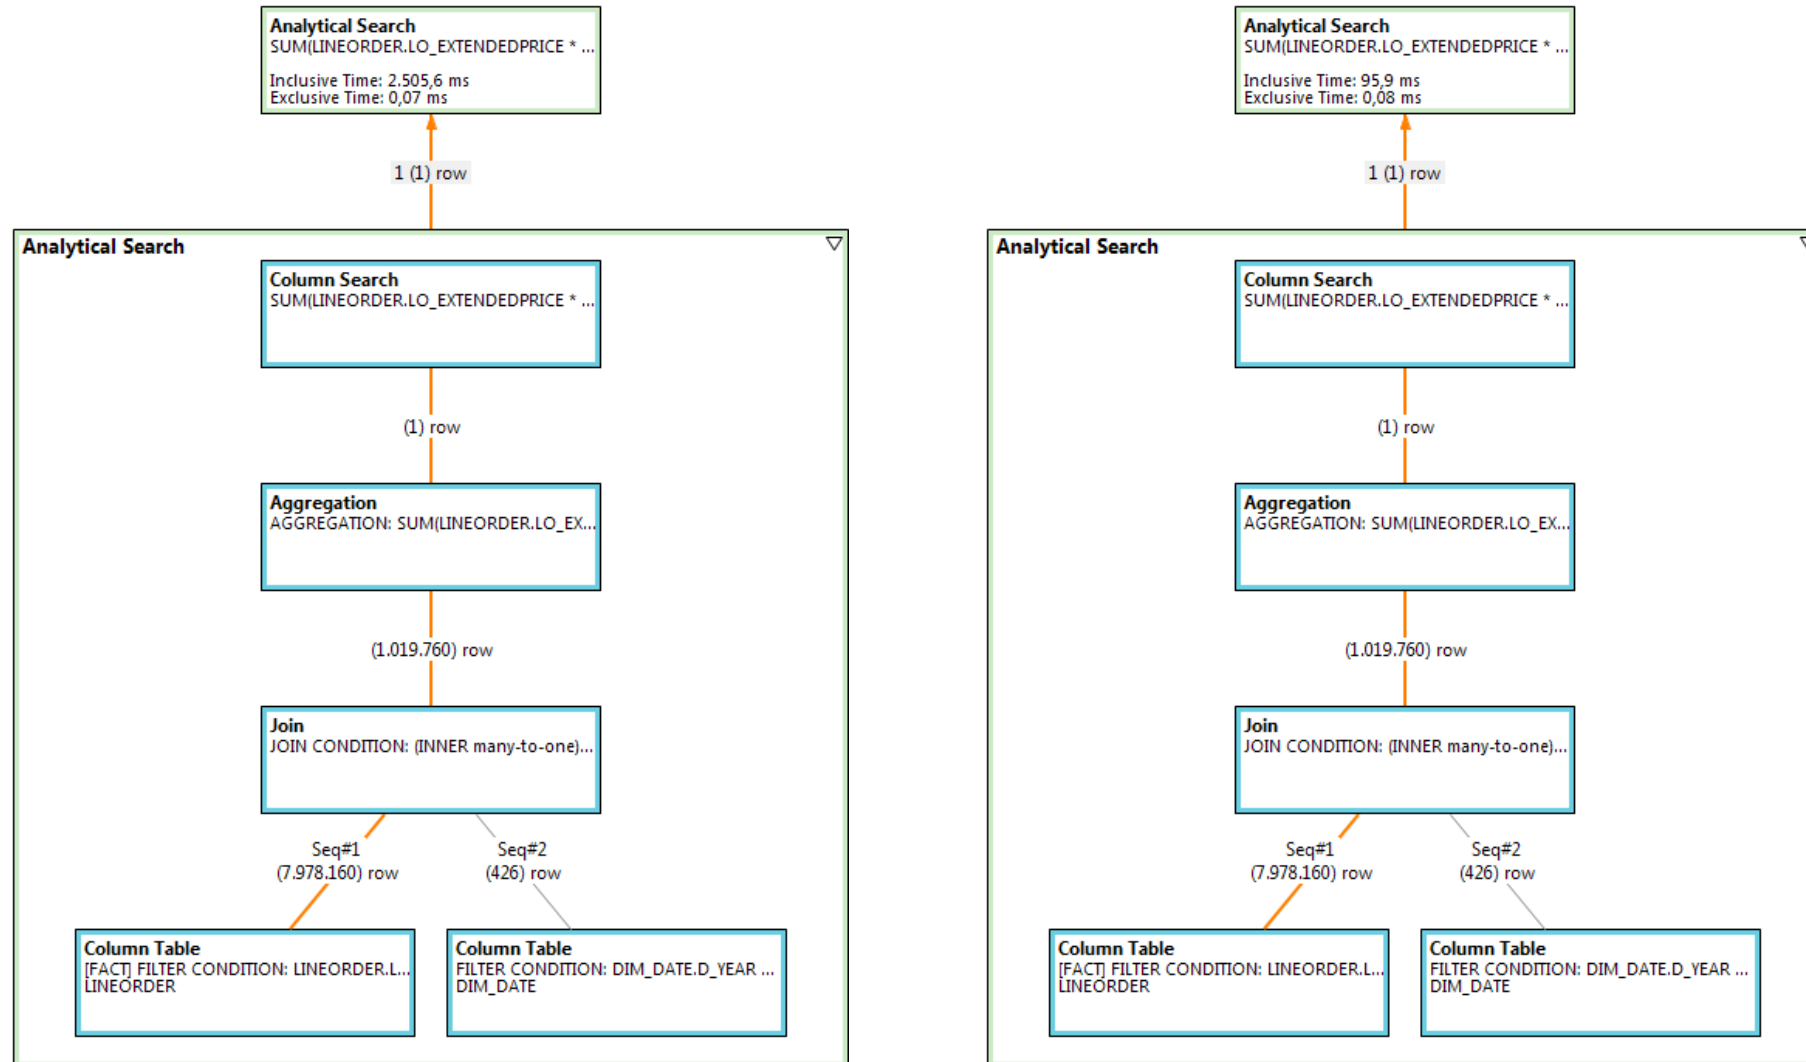


# Row Store mit und ohne Indizes

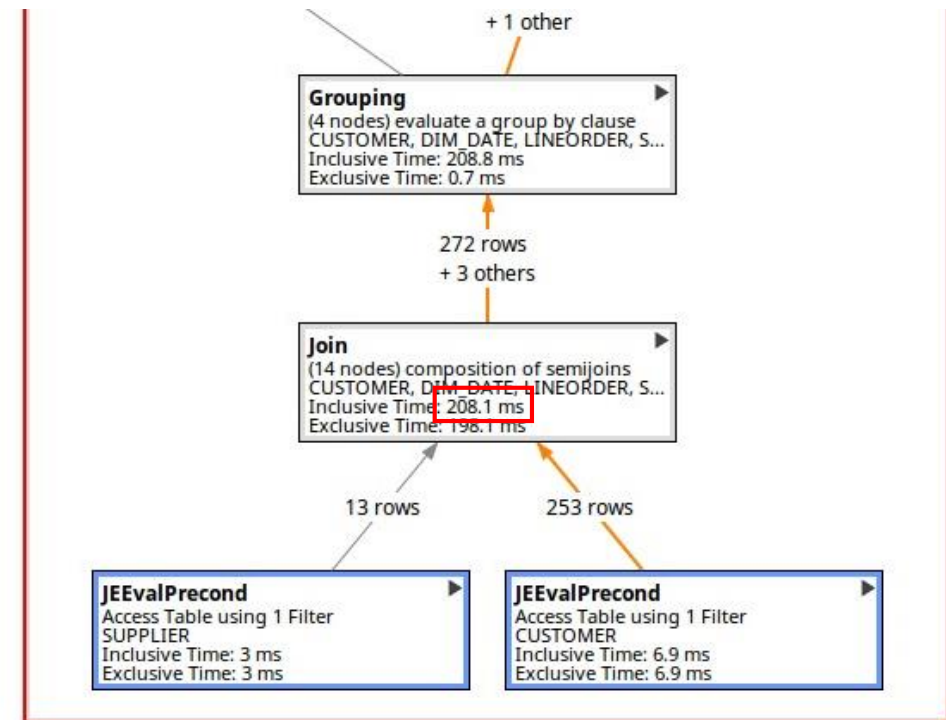
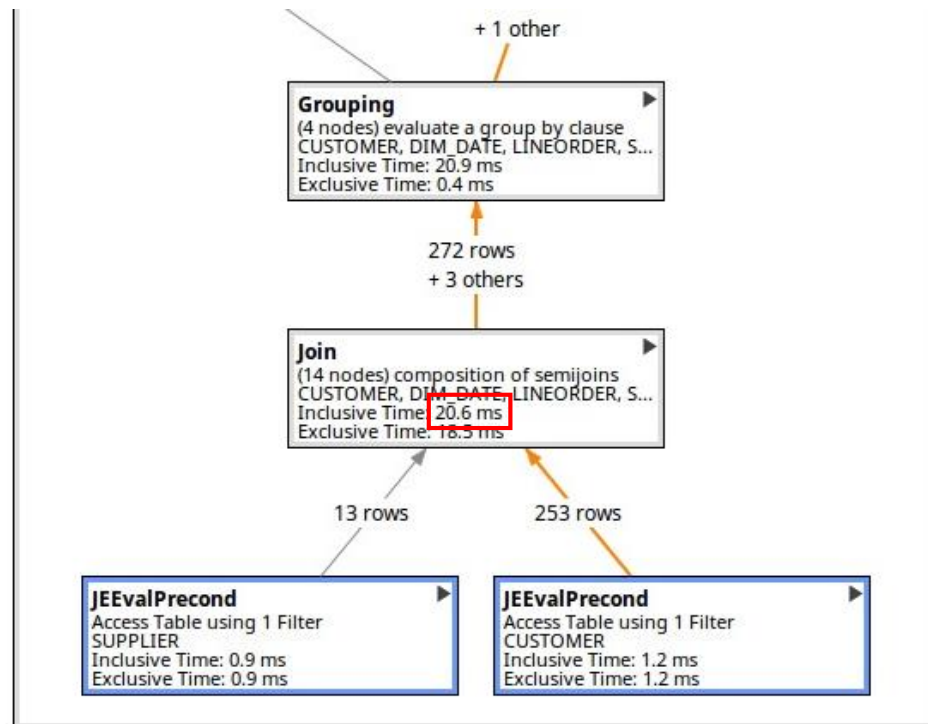





# Column Store ohne und mit Indizes



# Column Store mit und ohne Indizes (3.3)

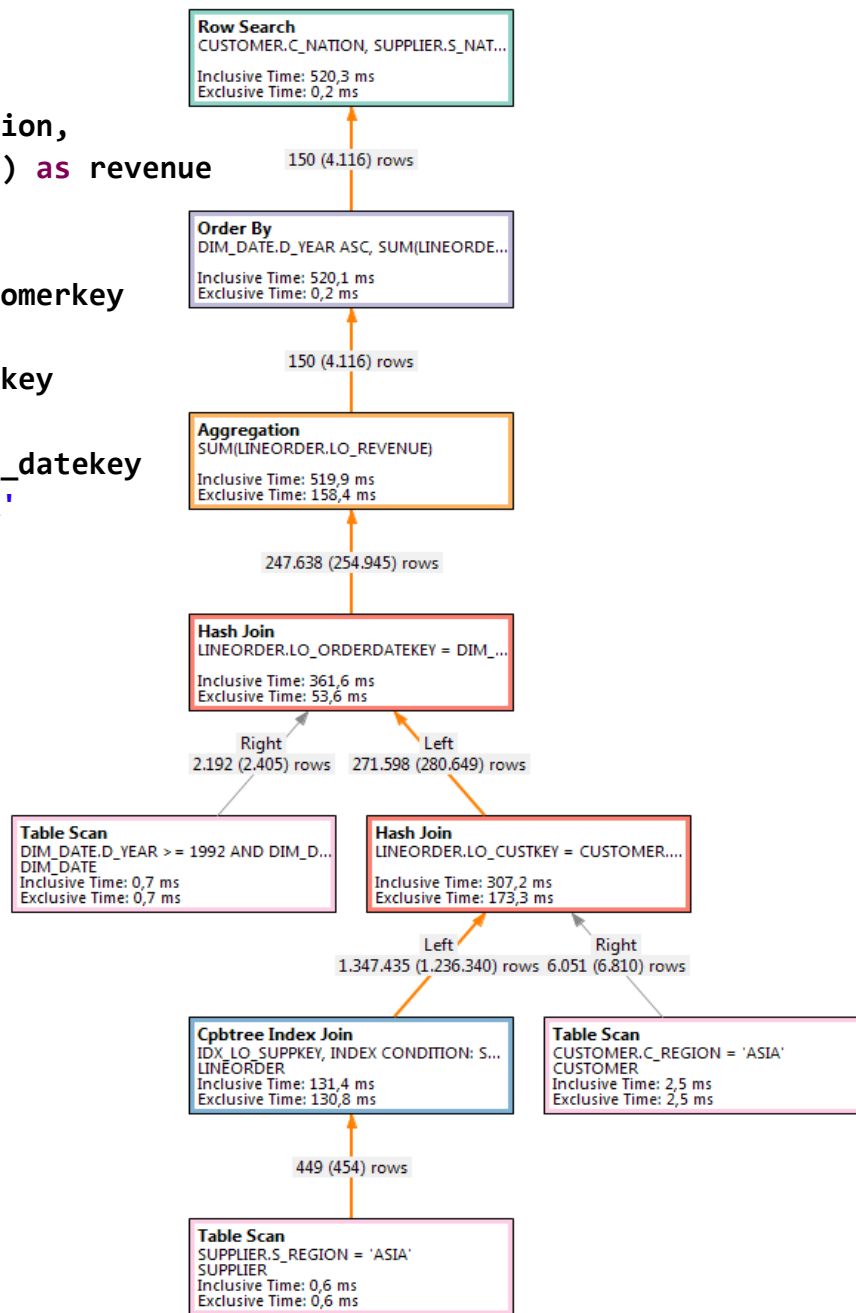


# Column vs. Row Store



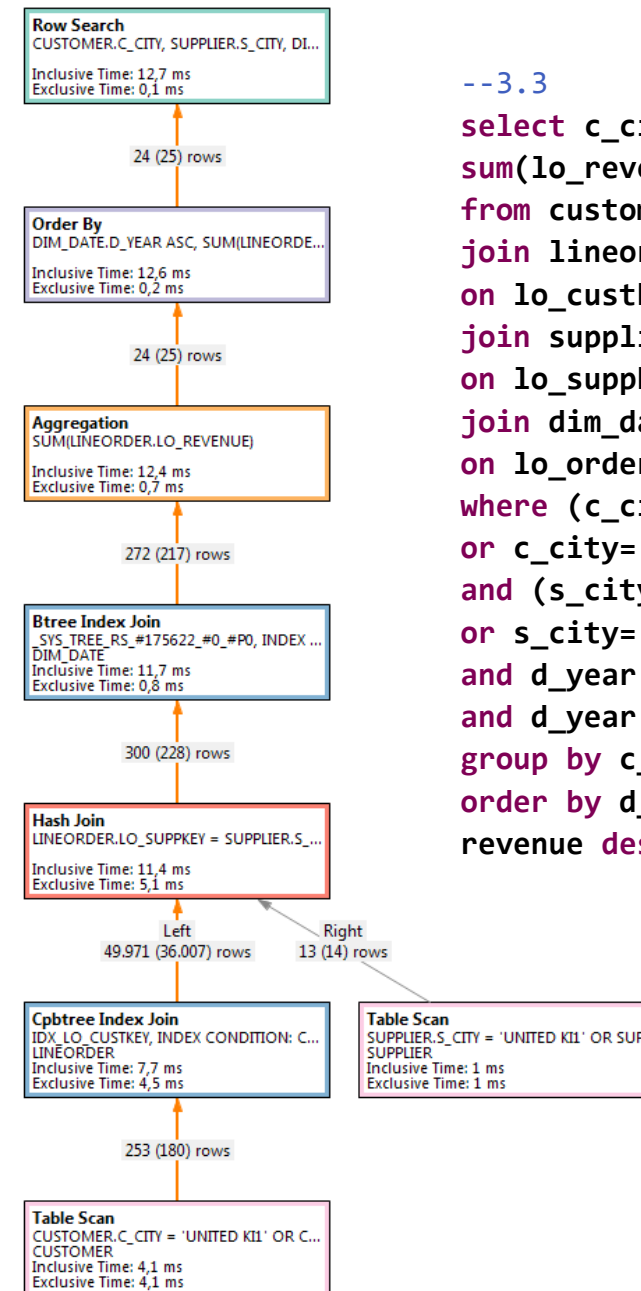
--3.1

```
select c_nation, s_nation,
d_year, sum(lo_revenue) as revenue
from customer
join lineorder
on lo_custkey = c_customerkey
join supplier
on lo_suppkey = s_suppkey
join dim_date
on lo_orderdatekey = d_datekey
where c_region = 'ASIA'
and s_region = 'ASIA'
and d_year >= 1992
and d_year <= 1997
group by c_nation,
s_nation, d_year
order by d_year asc,
revenue desc;
```



--3.3

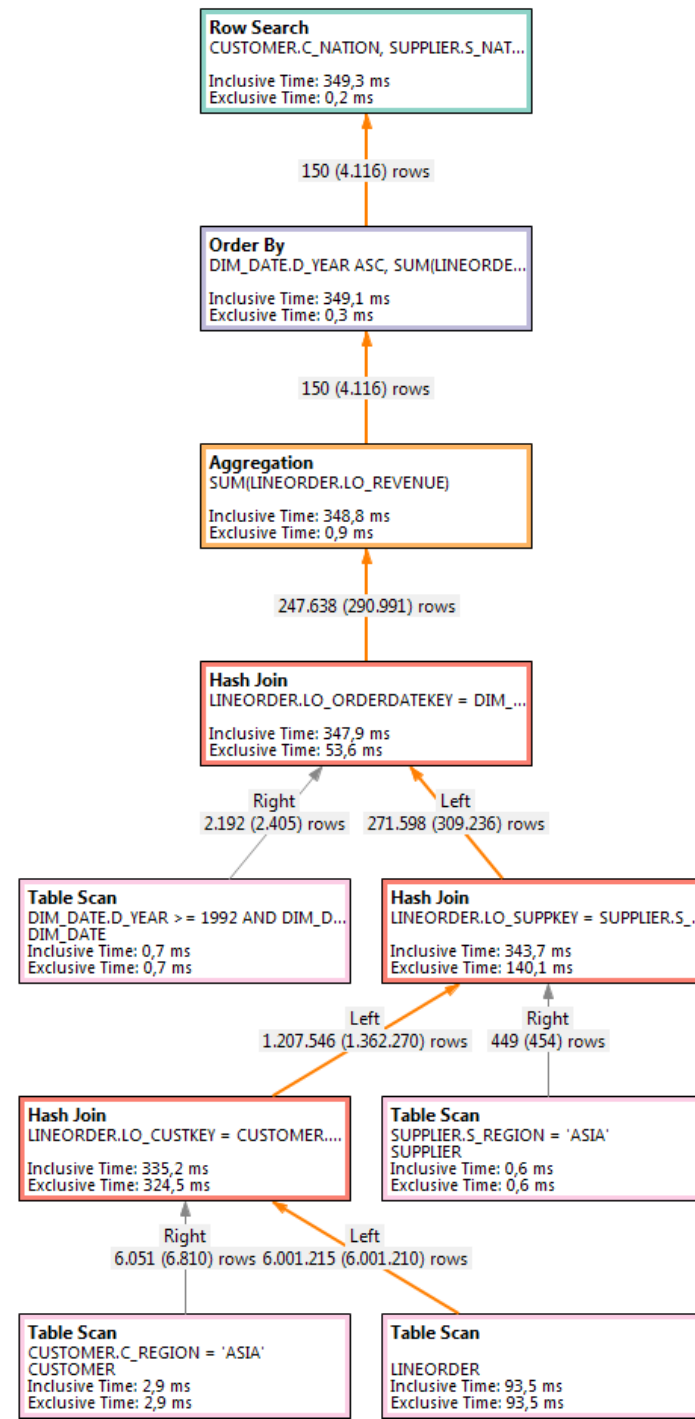
```
select c_city, s_city, d_year,
sum(lo_revenue) as revenue
from customer
join lineorder
on lo_custkey = c_customerkey
join supplier
on lo_suppkey = s_suppkey
join dim_date
on lo_orderdatekey = d_datekey
where (c_city='UNITED KI1'
or c_city='UNITED KI5')
and (s_city='UNITED KI1'
or s_city='UNITED KI5')
and d_year >= 1992
and d_year <= 1997
group by c_city, s_city, d_year
order by d_year asc,
revenue desc;
```



Q3.1

NO\_INDEX\_JOIN

1,653ms => 1,323ms




NO\_HASH\_JOIN => 3,609ms

RANGE\_JOIN => 1,696ms

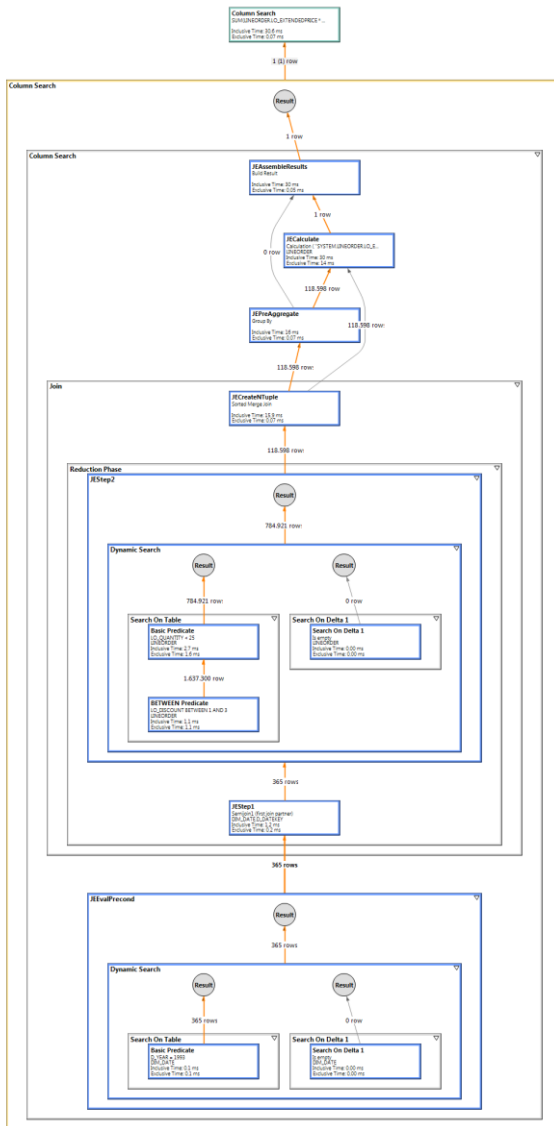
# Column vs. Row Store: Indizes

	Column Benchmark	Column Benchmark with Indizes	Row Benchmark	Row Benchmark with Index
Samples	10	10	10	10
Total	5783	5182	20325	18850
Max	1718	559	2382	1934
Min	442	505	1913	1871
Median	453	517	2001	1880
Average	578	518	2033	1885
StandardDeviation	380	16	121	17

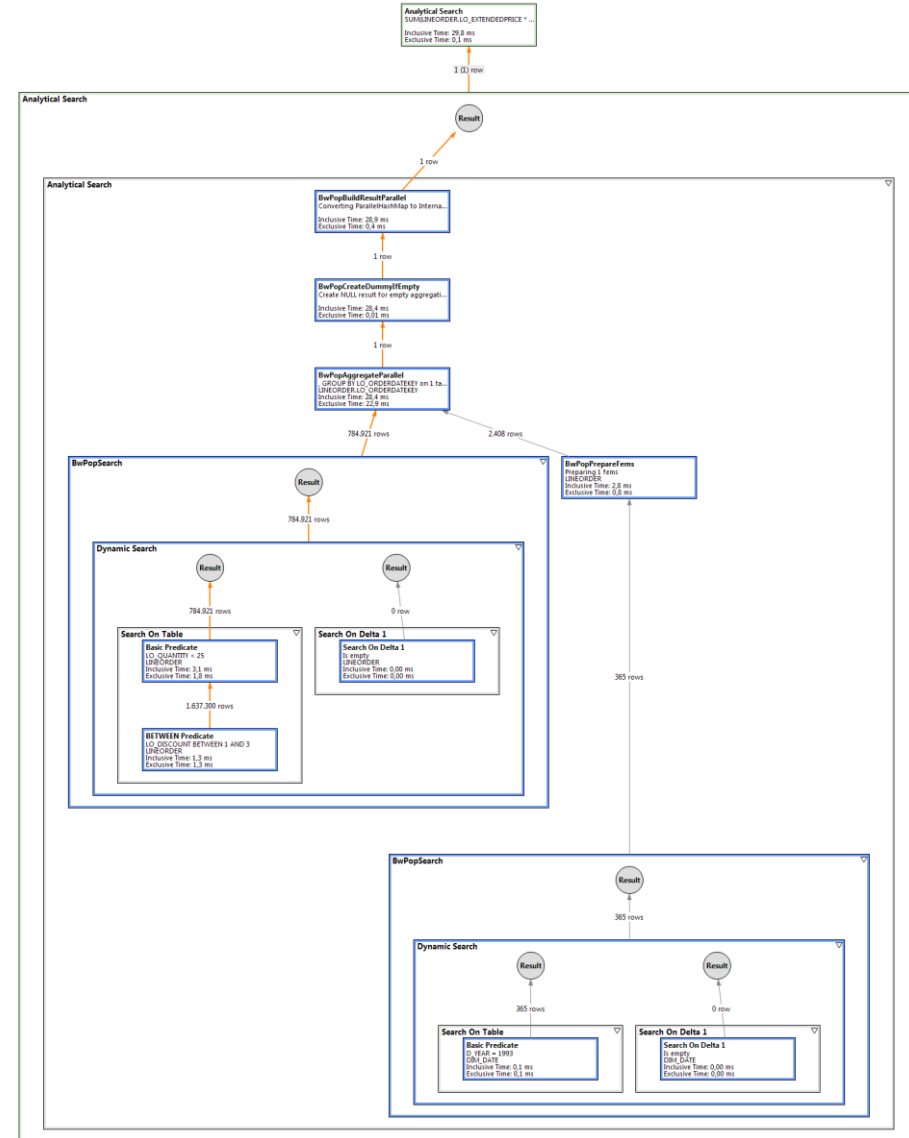
# Column vs. Row Store



## NO\_USE\_OLAP\_PLAN



## USE\_OLAP\_PLAN





A blue geometric graphic consisting of several thick lines that form a 3D cube-like structure. The lines are arranged in a way that creates a sense of depth and perspective, with some lines receding into the background and others coming forward. The graphic is centered on the page and partially obscured by a white horizontal band.

# Cube Präsentation

# SSBM Benchmark auf HANA

- Column Store ist schneller
- Indizes können Column Store in manchen Fällen beschleunigen
- Der Optimizer entscheidet ob OLAP Plan verwendet wird
- HANA braucht viel RAM

22.03.2018



# SSBM Benchmark auf HANA

- Column Store ist schneller
- Indizes können Column Store in manchen Fällen beschleunigen
- Der Optimizer entscheidet ob OLAP Plan verwendet wird
- HANA braucht viel RAM

22.03.2018

# DANKE!